



# Standardizing zkEVM acceleration API



Mamy Ratsimbazafy

ZK Engineering @ Taiko



mratsim



m\_ratsim

# Agenda

- ❖ The state of the PSE (Ethereum Foundation) zkEVM ecosystem
- ❖ The end goal
- ❖ The journey
- ❖ What's next?
- ❖ Collaborate

# The state of PSE (Ethereum Foundation) zkEVM ecosystem



# The users

- ❖ Co-processors:
  - Axiom
  - EZKL (zkML)
- ❖ Compilers, Languages & High-Level Frameworks
  - Lurk
  - Powdr
- ❖ Rollups:
  - Scroll
  - Taiko

# The open-source HW accel landscape

- ❖ On Halo2-KZG (Privacy Scaling Exploration)
  - <https://github.com/junyu0312/halo2> (March 2022)
  - <https://github.com/privacy-scaling-explorations/halo2/pull/79> (June 2022)
  - [https://github.com/superscalar-io/halo2\\_device\\_sample](https://github.com/superscalar-io/halo2_device_sample) (October 2023)
- ❖ On Halo2-IPA (Zcash)
  - <https://github.com/DelphinusLab/halo2-gpu-specific> (October 2023)
- ❖ On Arkworks
  - Icicle (Ingonyama)
  - Spparks (Supranational)

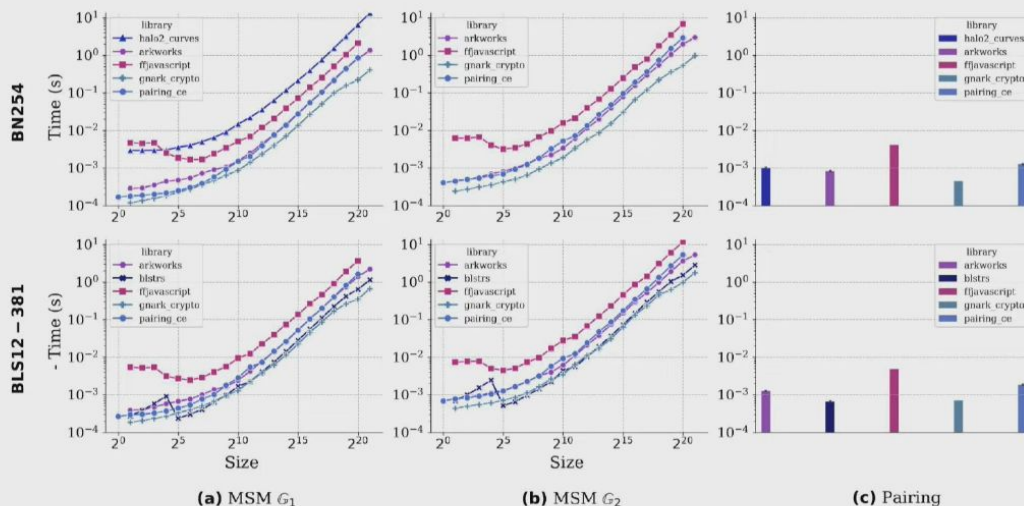
# The software landscape

- ❖ Towards state-of-the-art MSM / [halo2#187](#)
- ❖ <https://zka.lc/>
  - Benchmarking SNARKS, zkSummit 10  
<https://www.youtube.com/watch?v=fTGPUb07ebE>
  - <https://eprint.iacr.org/2023/1503>
  - Gnark is the fastest MSM provider benchmarked
- ❖ Current speed,  $2^{22}$  / 4M points, i9-11980HK, 8 cores:  
[PR halo2curves#86](#)
  - Halo2-KZG (PSE): 3.5s
  - Gnark (Consensus): 1.2s
  - Constantine (Taiko): 1.2s

# The software landscape

Ernstberger 2023 (zkSummit 10, Benchmarking SNARKS)

## MSM optimizations are dispersed!



# Software still has potential

- ❖ GPU renting costs
  - Competition with AI startups
- ❖ GPU embargos
- ❖ Parallelization on many-threaded CPUs
  - AMD EPYC 9654 96C/192T on 2 sockets hence 384 threads
  - Intel Xeon Platinum 8490H 60C/120T on 8 sockets hence 960 threads
- ❖ MSM parallelization level for the bucket method / Pippenger
  - MSM-level parallelism (partition points)
  - Window-level parallelism (partition scalar bits)
  - Bucket-level parallelism (need no collision when accumulating)



# Software future developments

- ❖ Specialized MSM, example [halo2#202](#) for bit-level keccak 40% speedup
- ❖ Prover-level parallelism, example EIP-4844 batch KZG commitments

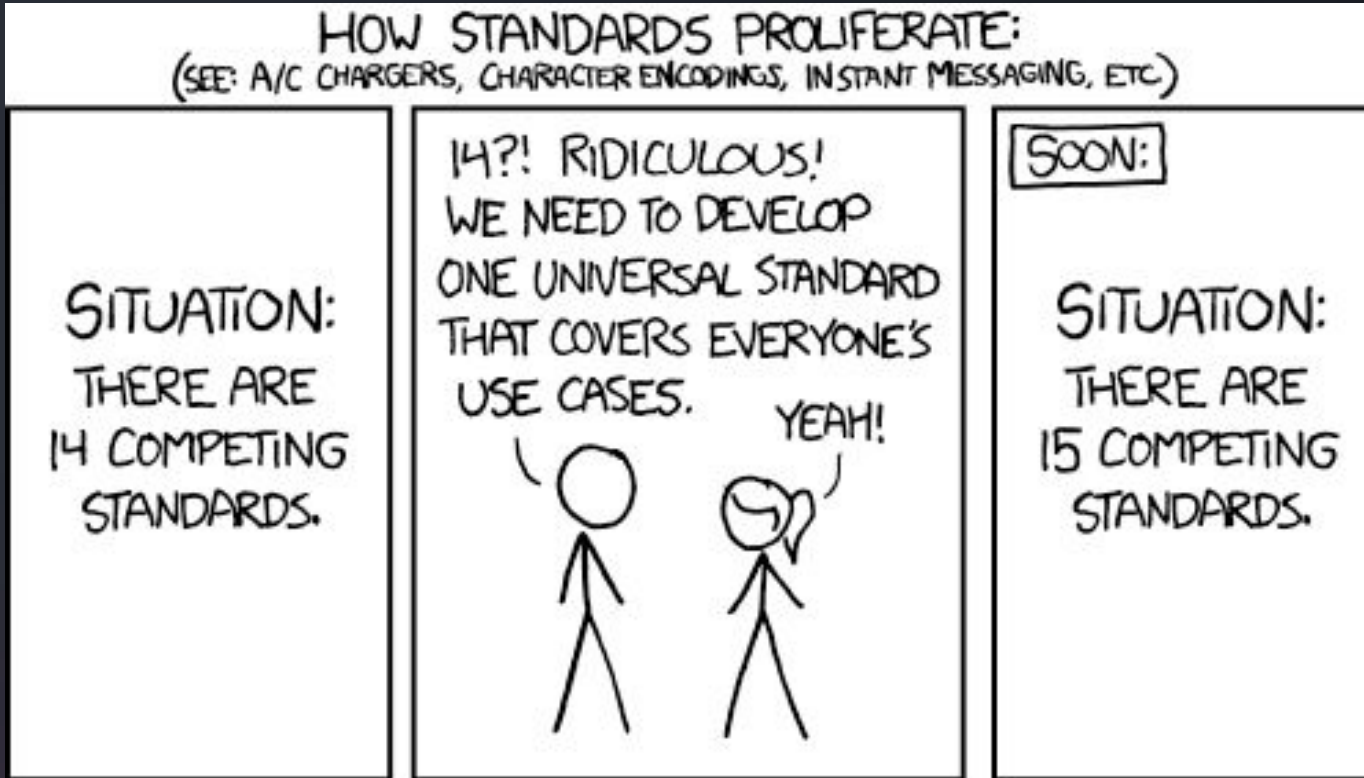
$$\begin{aligned} & e\left(\sum [r_i][proof_i]_1, [T]_2\right) \cdot \\ & e\left(\sum [r_i]([commitment_i]_1 - [eval\_at\_challenge_i]_1) \right. \\ & \quad \left. + \sum [r_i][z_i][proof_i]_1, [-1]_2\right) = 1 \end{aligned}$$

[https://github.com/mratsim/constantine/blob/5f7ba18f/constantine/commitments/kzg\\_polynomial\\_commitments\\_parallel.nim#L115](https://github.com/mratsim/constantine/blob/5f7ba18f/constantine/commitments/kzg_polynomial_commitments_parallel.nim#L115)

**The end goal**



# The end goal



# The end goal

- ❖ Disassociating proof systems from the computing backends
  - Provers are always evolving
  - The primitives they use MSMs, FFTs are not
- ❖ HW accel providers can serve many provers with a single codebase and small adapters instead of forking the complete codebase.
  - Less maintenance
  - Less bug surface
  - Focus on their specialty instead of having a proof-system person
- ❖ Proving libraries can focus on high-level proof-system, lookups optimizations
- ❖ Differential fuzzing and benchmarking
- ❖ End applications can pick the proving backend instead of being locked into the proving library default implementation

# The journey



# Standardizing the ABI

- ❖ Acceleration Abstraction Layer [Halo2#216](#)
- ❖ Proposes:
  - An encoding of field  $\mathbb{F}_r$  elements (64-bit word-endianess, limb-endianess, saturated limbs, Montgomery representation)
  - An encoding of elliptic curve  $G_1$  elements (homogeneous projective coordinates, default in Halo2-KZG / Ingonyama vs jacobian coordinates in Halo2-IPA / Arkworks / Spark ...)
  - An async API to issue multiple MSMs or FFTs on the accelerator in parallel
  - A context parameter to abstract all accelerator configuration from the prover library

# Standardizing the ABI

- ❖ Inspired by accelerators for image processing and machine learning
  - Nvidia Cuda and CuDNN
  - Intel oneDNN
  - OpenCV HAL (Hardware Abstraction Layer)
- ❖ Windows Kernel API
- ❖ An engine context
- ❖ An optional operation descriptor
  - We probably let the accel library handle memory allocation unlike what drivers or kernel do.

# Changes in Halo2

- ❖ ``best_fft`` appears twice
- ❖ ``best_multiexp`` appears 9 times
- ❖ All callers will need to pass a new ``engine`` context parameter
- ❖ End application will need to init the ``engine`` context with
  - Number of cores used (including reading env variables like `TAIKO_NUM_THREADS`)
  - GPU devices used



**What's next**



# Future improvements

- ❖ Async API for issuing multiple MSMs / FFTs in parallel (like EIP-4844 batch verification in Constantine)
- ❖ Caching abstraction via operation-specific descriptors / context
- ❖ Commitment-level Acceleration Layer
  - Instead of Arithmetic-level Acceleration Layer

**Collaborate**



# Collaborate

- ❖ <https://github.com/privacy-scaling-explorations/halo2/issues/216>
- ❖ [https://github.com/superscalar-io/halo2\\_device\\_sample](https://github.com/superscalar-io/halo2_device_sample)
- ❖ Working group on Telegram or Discord?
  - Hardware providers
  - Software libraries (Halo2-IPA, Halo2-KZG, Artworks, ...)
  - End-users (co-processors, languages & frameworks, rollups)

**The End**





# Standardizing zkEVM acceleration API



Mamy Ratsimbazafy

ZK Engineering @ Taiko



mratsim



m\_ratsim